

Multivariate Analysis Lecture 6: Sample Covariance Matrix and Wishart Distribution

Zhaoxia Yu
Professor, Department of Statistics

2023-04-20

Section 1

The Big Picture

The Big Picture: Univariate vs Multivariate

- **Review:** A random sample, denoted by X_1, \dots, X_n , from a (univariate) normal distribution $N(\mu, \sigma^2)$
 - What are the distributions of \bar{X}, s^2 ? What useful statistics can be constructed?
- **New material:** A random sample, denoted by $\mathbf{X}_1, \dots, \mathbf{X}_n$, from a multivariate normal distribution $N(\mu, \Sigma)$
 - What are the distributions of $\bar{\mathbf{X}}, \mathbf{S}$? What useful statistics can be constructed?

The Big Picture: Univariate

- A random sample, denoted by X_1, \dots, X_n , from a (univariate) normal distribution $N(\mu, \sigma^2)$
- Let $\mathbf{X}_{n \times 1} = (X_1, \dots, X_n)^T$. It is random vector with a multivariate normal distribution, i.e.,

$$\mathbf{X}_{n \times 1} = (X_1, \dots, X_n)^T \sim \mathbf{N}(\mu \mathbf{1}, \sigma^2 \mathbf{I})$$

- 1 $\bar{X} \sim N(\mu, \sigma^2/n)$
- 2 $\frac{(n-1)s^2}{\sigma^2} \sim \chi_{n-1}^2$
- 3 Independence between \bar{X} and s^2 .
- 4 a t-statistic is

$$\frac{\frac{\bar{X} - \mu}{\sqrt{\sigma^2/n}}}{\sqrt{\frac{(n-1)s^2/\sigma^2}{n-1}}} = \frac{\sqrt{n}(\bar{X} - \mu)}{s}$$

It follows the t-distribution with $n-1$ degrees of freedom, denoted by t_{n-1} .

The Big Picture: Multivariate

- A random sample $\mathbf{X}_1, \dots, \mathbf{X}_n$ from a multivariate normal distribution $\mathbf{N}(\boldsymbol{\mu}, \boldsymbol{\Sigma})$.
- Let

$$\mathbf{X}_{n \times p} = \begin{pmatrix} \mathbf{X}_1^T \\ \vdots \\ \mathbf{X}_n^T \end{pmatrix}$$

\mathbf{X} follows a matrix normal distribution.

- 1 Sample mean vector follows a multivariate normal, i.e., $\bar{\mathbf{X}} \sim \mathbf{N}(\boldsymbol{\mu}, \boldsymbol{\Sigma}/n)$
- 2 Sample covariance matrix $(n-1)\mathbf{S}$ follows a Wishart distribution, i.e., $(n-1)\mathbf{S} \sim \text{Wishart}_p(n-1, \boldsymbol{\Sigma})$
- 3 Independence between $\bar{\mathbf{X}}$ and \mathbf{S} .
- 4 Hotelling's T^2 : $T^2 = (\bar{\mathbf{X}} - \boldsymbol{\mu})^T \left(\frac{\mathbf{S}}{n}\right)^{-1} (\bar{\mathbf{X}} - \boldsymbol{\mu})$

The Big Picture: outline

- Sample variance and chi-squared distribution
- Sample covariance matrix and Wishart distribution
- Hotelling's T^2
- Maximum likelihood estimate

Section 2

Sample Variance

Sample Variance and Chi-squared Distribution

- Let $\mathbf{X} = (X_1, \dots, X_n)$ denote a random sample from $N(\mu, \sigma^2)$.
- Equivalently, $\mathbf{X} \sim N(\mu \mathbf{1}, \sigma^2 \mathbf{I})$.
- Let $s^2 = \frac{1}{n-1} \sum_{i=1}^n (X_i - \bar{X})^2$ denote the sample variance.
- We would like to show that

$$\frac{(n-1)s^2}{\sigma^2} \sim \chi_{n-1}^2$$

- Outline of proof
 - 1 Projection matrices
 - 2 Chi-squared distribution
 - 3 Rewrite $(n-1)s^2/\sigma^2$ as the sum of squared $N(0, 1)$ random variables

Projection Matrices

- A projection matrix is a square matrix that is both idempotent and symmetric

$$\mathbf{P}^2 = \mathbf{P}, \mathbf{P} = \mathbf{P}^T$$

Projection Matrices

- Suppose \mathbf{P} is a projection matrix. We have
 - The eigenvalues of \mathbf{P} has eigenvalues are either 0 or 1, and the number of 1's is the same as the rank of the projection matrix.
 - $\text{tr}(\mathbf{P}) = \text{rank}(\mathbf{P})$
 - The spectral decomposition of \mathbf{P} is

$$\mathbf{P} = \sum_{i=1}^r \gamma_i \gamma_i^T$$

where $r = \text{rank}(\mathbf{P})$, and $(\gamma_1, \dots, \gamma_r)$ are orthogonal vectors of norm 1, i.e.,

$$\gamma_i^T \gamma_j = \begin{cases} 1 & \text{if } i = j \\ 0 & \text{if } i \neq j \end{cases}$$

A Special Projection Matrix: the Centering Matrix

- The centering matrix $\mathbb{C} = \mathbf{I} - \frac{1}{n}\mathbf{1}\mathbf{1}^T$ is a very special matrix.
- It is a projection matrix, which is defined as both symmetric and idempotent:
 - $\mathbb{C}^T = \mathbb{C}$ (symmetric)
 - $\mathbb{C}^2 = \mathbb{C}$ (idempotent)
- One important result about a projection matrix is that its eigenvalues are either zero or one.
- By properties of projection matrices, we have
 - $\text{rank}(\mathbb{C}) = \text{tr}(\mathbb{C}) = n - 1$
 - $\mathbb{C} = \sum_{j=1}^{n-1} \gamma_j \gamma_j^T$

A Special Projection Matrix: the Centering Matrix

- The centering matrix centers data
- Univariate: Let $\mathbf{X}_{n \times 1}$ be a random sample from $N(\mu, \sigma^2)$, i.e.,

$$\mathbf{X}_{n \times 1} \sim N(\mu \mathbf{1}, \sigma^2 \mathbf{I})$$

$\mathbb{C}\mathbf{X}$ is a linear function of \mathbf{X} and it can be verified that $\mathbb{C}\mathbf{1} = \mathbf{0}$, we have

$$E[\mathbb{C}\mathbf{X}] = \mu \mathbb{C}\mathbf{1} = \mathbf{0}$$

- Multivariate: Let $\mathbf{X}_{n \times p}$ be a random sample from $N(\mu, \Sigma)$
Similarly, it can be shown that $\mathbb{C}\mathbf{X}$ has mean $\mathbf{0}_{n \times p}$. We have verified this numerically.
- In either situation, we have $\mathbb{C}\mathbf{X} = \mathbb{C}(\mathbf{X} - E[\mathbf{X}])$ This fact will be used later.

Chi-squared distribution

- **Definition.** Let Z_1, Z_2, \dots, Z_k be independent standard normal random variables. Then, the sum of squares $Q = Z_1^2 + Z_2^2 + \dots + Z_k^2$ has a chi-squared distribution with k degrees of freedom, denoted by χ_k^2 .
- **Alternative definition.** Let $\mathbf{Z}_{k \times 1} \sim N(\mathbf{0}, \mathbf{I})$. We say $\|\mathbf{Z}\|^2 = \mathbf{Z}^T \mathbf{Z}$ follows χ_k^2 .
- The PDF of a chi-squared random variable with k degrees of freedom is given by:

$$f(x) = \frac{1}{2^{k/2} \Gamma(k/2)} x^{k/2-1} e^{-x/2}, \quad x > 0$$

where $\Gamma(\cdot)$ is the gamma function.

Chi-squared distribution

- The chi-squared distribution is a special case of the gamma distribution, where the shape parameter is $k/2$ and the rate parameter is $1/2$.
- The MGF of a chi-squared random variable with k degrees of freedom is:

$$M_X(t) = (1 - 2t)^{-k/2}$$

- The mean and variance of a chi-squared random variable with k degrees of freedom are:

$$E[X] = k, \text{Var}[X] = 2k$$

Construct Chi-squared R.V.s using Normal R.V.s and Projection Matrices

- Let $\mathbf{P}_{n \times n}$ be a projection matrix with rank r and let $\mathbf{Z}_{n \times 1} \sim N(\mathbf{0}, \mathbf{I})$

$$\begin{aligned}\mathbf{Z}^T \mathbf{P} \mathbf{Z} &= \mathbf{Z}^T \sum_{i=1}^r \gamma_i \gamma_i^T \mathbf{Z} = \sum_{i=1}^r \mathbf{Z}^T \gamma_i \gamma_i^T \mathbf{Z} \\ &= \sum_{i=1}^r (\gamma_i^T \mathbf{Z})^T (\gamma_i^T \mathbf{Z})\end{aligned}$$

Let $Y_i = \gamma_i^T \mathbf{Z}$. Note that Y_i is univariate and it is a linear combination of \mathbf{Z} , from which we can show that $Y_i \sim N(0, 1)$.

- Note that $\mathbf{Z}^T \mathbf{P} \mathbf{Z} = \sum_{i=1}^r Y_i^2$. By the definition of chi-squared distribution, we have $\mathbf{Z}^T \mathbf{P} \mathbf{Z} \sim \chi_r^2$

The Sample Variance

- We have shown that
 - $\mathbb{C} = \mathbf{I} - \frac{1}{n}\mathbf{1}\mathbf{1}^T$
 - $\mathbb{C}^T = \mathbb{C}$, $\mathbb{C}^2 = \mathbb{C}$.
 - It is a projection matrix with rank $n - 1$ and

$$\mathbb{C} = \sum_{j=1}^{n-1} \gamma_j \gamma_j^T$$

-The centering matrix does center data, i.e.,

$$\mathbb{C}\mathbf{X} = \mathbb{C}(\mathbf{X} - E[\mathbf{X}])$$

- $(n - 1)s^2 = \mathbf{X}^T \mathbb{C} \mathbf{X}$, where

The Sample Variance

- Therefore,

$$\begin{aligned}\frac{(n-1)s^2}{\sigma^2} &= \frac{(\mathbf{X} - E[\mathbf{X}])^T}{\sigma} \mathbb{C}^T \mathbb{C} \frac{(\mathbf{X} - E[\mathbf{X}])}{\sigma} \\ &= \frac{(\mathbf{X} - E[\mathbf{X}])^T}{\sigma} \mathbb{C} \frac{(\mathbf{X} - E[\mathbf{X}])}{\sigma}\end{aligned}$$

The Sample Variance

- Let

$$\mathbf{Z} = \frac{(\mathbf{X} - E[\mathbf{X}])}{\sigma}$$

- Easy to see that $\mathbf{Z} \sim N(\mathbf{0}, \mathbf{I})$. Thus,

$$\frac{(n-1)s^2}{\sigma^2} = \mathbf{Z}^T \mathbb{C} \mathbf{Z}$$

- Use the result in previous slides, we have

$$\frac{(n-1)s^2}{\sigma^2} = \mathbf{Z}^T \mathbb{C} \mathbf{Z} \sim \chi_{n-1}^2$$

Section 3

Sample Covariance

The Sample Covariance from A MVN Random Sample

- Let $\mathbf{X}_1, \dots, \mathbf{X}_n \stackrel{iid}{\sim} N(\boldsymbol{\mu}, \boldsymbol{\Sigma})$.
- Recall that the sample covariance matrix is defined as

$$\mathbf{S} = \frac{1}{n-1} \sum_{i=1}^n (\mathbf{X}_i - \bar{\mathbf{X}})(\mathbf{X}_i - \bar{\mathbf{X}})^T$$

- We have shown that

$$(n-1)\mathbf{S} = \mathbf{X}^T \mathbb{C} \mathbf{X}$$

where \mathbf{X} is the $n \times p$ random matrix.

The Sample Covariance from A MVN Random Sample

- The goal is to show that $(n - 1)\mathbf{S}$ follows a Wishart distribution. More specifically, we would like to show that

$$(n - 1)\mathbf{S} \sim \text{Wishart}_p(n - 1, \Sigma)$$

- Outline of proof
 - 1 Wishart-distribution
 - 2 Rewrite $(n - 1)\mathbf{S}$
 - 3 Apply properties of a projection matrix
 - 4 Use the definition of Wishart distribution

Wishart Distribution

- The Wishart distribution is named after the British statistician John Wishart, who introduced it in his 1928 paper published in Biometrika.
- Wishart was interested in the problem of estimating the covariance matrix of a multivariate normal distribution.
- Wishart showed that the sample covariance matrix follows a particular probability distribution that we now call the Wishart distribution.
- The Wishart distribution has become a fundamental tool in multivariate statistical analysis

Definition of Wishart Distribution

- A Wishart distribution can be defined in the following way
- Let \mathbf{W} be a $p \times p$ random matrix. We say \mathbf{W} follows $Wishart_p(k, \Sigma)$ if \mathbf{W} can be written as $\mathbf{W} = \mathbf{X}^T \mathbf{X}$ where \mathbf{X} denotes the random matrix formed by a random sample of size k from MVN $N(\mathbf{0}, \Sigma)$.
- The definition indicates that if we have a random sample $\mathbf{X}_1, \dots, \mathbf{X}_k$ from $N(\mathbf{0}, \Sigma)$, then $\mathbf{X}^T \mathbf{X} = \sum_{i=1}^k \mathbf{X}_i \mathbf{X}_i^T \sim Wishart_p(k, \Sigma)$.
- Remark: $E[\mathbf{W}] = k\Sigma$.

Wishart vs Chi-squared

- **Wishart:** If $\mathbf{X}_1, \dots, \mathbf{X}_k \stackrel{iid}{\sim} N(\mathbf{0}, \boldsymbol{\Sigma})$, then

$$\mathbf{X}^T \mathbf{X} = \sum_{i=1}^k \mathbf{X}_i \mathbf{X}_i^T \sim \text{Wishart}_p(k, \boldsymbol{\Sigma}), \text{ where } \mathbf{X}_{k \times p} = \begin{pmatrix} X_1^T \\ \vdots \\ X_k^T \end{pmatrix}$$

- **Chi-squared:** If $X_1, \dots, X_k \stackrel{iid}{\sim} N(0, 1)$, then

$$\mathbf{X}^T \mathbf{X} = \sum_{i=1}^k X_i^2 \sim \chi_k^2, \text{ where } \mathbf{X}_{k \times 1} = \begin{pmatrix} X_1 \\ \vdots \\ X_k \end{pmatrix}$$

Wishart vs Chi-squared (continued)

- When $p = 1$,

$$W = \sum_{i=1}^k X_i^2 = \sigma^2 \sum_{i=1}^k \left(\frac{X_i}{\sigma} \right)^2 \sim \sigma^2 \chi_k^2$$

The Sample Covariance Matrix

- Let $\mathbf{X}_1, \dots, \mathbf{X}_n$ be a random sample from $N(\boldsymbol{\mu}, \boldsymbol{\Sigma})$. The $\mathbf{X}_{n \times p}$ follows a matrix normal distribution:

$$\mathbf{X} \sim N(\mathbf{1}_n \otimes \boldsymbol{\mu}^T, \boldsymbol{\Sigma}, \mathbf{I}_n)$$

- The sample covariance $(n-1)\mathbf{S} = \mathbf{X}^T \mathbb{C} \mathbf{X}$ is based on the centered data. The definition of Wishart distribution is not applicable immediately.
- Next we show that $(n-1)\mathbf{S}$ follows $Wishart_p(n-1, \boldsymbol{\Sigma})$.

The Sample Covariance Matrix

- Rewrite $(n - 1)\mathbf{S}$:

$$\begin{aligned}(n - 1)\mathbf{S} &= \mathbf{X}^T \mathbb{C}^T \mathbb{C} \mathbf{X} = (\mathbb{C}\mathbf{X})^T (\mathbb{C}\mathbf{X}) \\&= (\mathbb{C}\mathbf{X})^T \mathbb{C} (\mathbb{C}\mathbf{X}) \\&= (\mathbb{C}\mathbf{X})^T \sum_{j=1}^{n-1} \gamma_j \gamma_j^T (\mathbb{C}\mathbf{X}) \\&= \sum_{j=1}^{n-1} (\gamma_j^T \mathbb{C}\mathbf{X})^T (\gamma_j^T \mathbb{C}\mathbf{X})\end{aligned}$$

The Sample Covariance Matrix

- Let $Y_i = (\gamma_i^T \mathbb{C} \mathbf{X})^T$, we have
 - $E[Y_i] = 0$ because \mathbb{C} is the centering matrix
 - In the following, we show that Y_i and Y_j are uncorrelated for $i \neq j$:

$$\begin{aligned}\text{Cov}[Y_i, Y_j] &= E[(Y_i - \mathbf{0})(Y_j - \mathbf{0})^T] \\ &= E[Y_i Y_j^T] \\ &= E[(\gamma_i^T \mathbb{C} \mathbf{X})^T (\gamma_j^T \mathbb{C} \mathbf{X})] \\ &= E[\mathbf{X}^T \mathbb{C} \gamma_i \gamma_j^T \mathbb{C} \mathbf{X}] \\ &= \mathbf{0}\end{aligned}$$

The last step is true because for $i \neq j$, $\gamma_i \gamma_j^T = 0$

The Sample Covariance Matrix

Since Y_i and Y_j are two linear combinations of the same MVN distributed random matrix (or its vectorized version), we have Y_i and Y_j are independent for $i \neq j$.

- It can also be shown that $Y_i \sim N(\mathbf{0}, \Sigma)$.
- By the definition of Wishart, we can conclude that

$$(n - 1)\mathbf{S} \sim \text{Wishart}_p(n - 1, \Sigma)$$

Some Interesting Results

- Consider a random sample from MVN $N(\boldsymbol{\mu}, \boldsymbol{\Sigma})$. Let \mathbf{S} denote the sample covariance matrix.
- We have already shown that $(n-1)\mathbf{S} \sim \text{Wishart}_p(n-1, \boldsymbol{\Sigma})$
- What is the distribution of a diagonal element of $(n-1)\mathbf{S}$?
- What is the distribution of the sum of elements of $(n-1)\mathbf{S}$?
- What is the distribution of $(n-1)\mathbf{BSB}^T$ where B is a fixed $q \times p$ matrix?
- If time permits, we will run some simulations

Some Interesting Results (continued)

- If you cannot get the answer to the last question, let's use the definition of Wishart distribution.
- Let $\mathbf{W} = (n - 1)S$. Because it follows $Wishart_p(n - 1, \mathbf{\Sigma})$, we know that $\mathbf{W} = \sum_{j=1}^{n-1} \mathbf{Z}_j \mathbf{Z}_j^T$ where \mathbf{Z}_j 's are iid from $N(\mathbf{0}, \mathbf{\Sigma})$.
- Then

$$\begin{aligned}(n - 1)\mathbf{B}S\mathbf{B}^T &= \mathbf{B} \sum_{j=1}^{n-1} \mathbf{Z}_j \mathbf{Z}_j^T \mathbf{B} = \sum_{j=1}^{n-1} \mathbf{B} \mathbf{Z}_j \mathbf{Z}_j^T \mathbf{B}^T \\ &= \sum_{j=1}^{n-1} (\mathbf{B} \mathbf{Z}_j)(\mathbf{B} \mathbf{Z}_j)^T\end{aligned}$$

Some Interesting Results (continued)

Let $\mathbf{Y}_j = \mathbf{B}\mathbf{Z}_j$. Note that it is a linear function of \mathbf{Z}_j ; therefore

$$\mathbf{Y}_j \sim N(\mathbf{0}, \mathbf{B}\Sigma\mathbf{B}^T)$$

and the \mathbf{Y}_j 's are iid (because ...).

By the definition of Wishart distribution, we have

$$(n-1)\mathbf{BSB}^T \sim \text{Wishart}_q(n-1, \mathbf{B}\Sigma\mathbf{B}^T)$$

Section 4

Hotelling's T^2

The Hotelling's T^2 Statistic

- Finally we are ready to introduce Hotelling's
- The student's t is used for making inference of mean(s) of normal distribution(s)
- Hotelling generalized the student's t , which is for univariate, to Hotelling's T^2 , which is the multivariate version

Definition Hotelling's T^2

- **Definition.** We say a random variable follows Hotelling's T^2 $T_{p,\nu}^2$ if the random variable can be written as $\mathbf{Z}^T \left(\frac{\mathbf{W}}{\nu} \right)^{-1} \mathbf{Z}$ where
 - 1 $\mathbf{Z} \sim N(\mathbf{0}, \Sigma)$
 - 2 $\mathbf{W} \sim W_p(\nu, \sigma)$
 - 3 $\mathbf{Z} \perp \mathbf{W}$

One-Sample Hotelling T^2

- Let $\mathbf{X}_1, \mathbf{X}_2, \dots, \mathbf{X}_n$ be a random sample from a multivariate normal distribution with mean vector $\boldsymbol{\mu}$ and covariance matrix $\boldsymbol{\Sigma}$.
- The sample mean vector and sample covariance matrix are denoted by $\bar{\mathbf{X}}$ and \mathbf{S} , respectively.
- The null hypothesis of interest $H_0 : \boldsymbol{\mu} = \boldsymbol{\mu}_0$
- The one-sample Hotelling T^2 is defined as

$$T^2 = (\hat{\boldsymbol{\mu}} - \boldsymbol{\mu}_0)^T (\text{Cov}(\hat{\boldsymbol{\mu}}))^{-1} (\hat{\boldsymbol{\mu}} - \boldsymbol{\mu}_0)$$

One-Sample Hotelling T^2 (continued)

- To see that T^2 does follow Hotelling's T^2 , we rewrite it

$$\begin{aligned} T^2 &= (\hat{\mu} - \mu_0)^T (\text{Cov}(\hat{\mu}))^{-1} (\hat{\mu} - \mu_0) \\ &= (\bar{\mathbf{X}} - \mu_0)^T (\text{Cov}(\bar{\mathbf{X}}))^{-1} (\bar{\mathbf{X}} - \mu_0) \\ &= (\bar{\mathbf{X}} - \mu_0)^T \left(\frac{S}{n} \right)^{-1} (\bar{\mathbf{X}} - \mu_0) \\ &= [\sqrt{n}(\bar{\mathbf{X}} - \mu_0)]^T \left(\frac{(n-1)S}{n-1} \right)^{-1} [\sqrt{n}(\bar{\mathbf{X}} - \mu_0)] \end{aligned}$$

- We have shown that all the three conditions for constructing a Hotelling's T^2 are satisfied
- As a result, $T^2 \sim T^2_{p, n-1}$.

Hotelling's T^2 Distribution vs F Distribution

Hotelling's T^2

Claim: $T_{p,\nu}^2 \sim \frac{\nu p}{\nu+1-p} F_{p,\nu+1-p}$.

For the T^2 statistic, we have $T^2 \stackrel{H_0}{\sim} \frac{(n-1)p}{n-p} F_{p,n-p}$. We reject H_0 at significance level α when $T^2 > \frac{(n-1)p}{n-p} F_{p,n-p,1-\alpha}$.

Corollary.

$$\frac{n-p}{p} (\bar{X} - \mu_0)^T (\hat{\Sigma})^{-1} (\bar{X} - \mu_0) \stackrel{H_0}{\sim} F_{p,n-p}$$

where $\hat{\Sigma} = \frac{1}{n} X^T H X = \frac{1}{n} \sum_{i=1}^n (X_i - \bar{X})(X_i - \bar{X})^T = \frac{(n-1)S}{n}$.

Section 5

MLE

MLE: Introduction

- The maximum likelihood estimate (MLE) is a widely used method for estimating the parameters of a statistical model.
- In this presentation, we will focus on the MLE for a multivariate normal distribution.

MLE: Multivariate Normal Distribution

- A random vector \mathbf{X} follows a p -dimensional multivariate normal distribution with mean vector $\boldsymbol{\mu}$ and covariance matrix $\boldsymbol{\Sigma}$, denoted by $\mathbf{X} \sim \mathcal{N}_p(\boldsymbol{\mu}, \boldsymbol{\Sigma})$, if its probability density function is given by:

$$f_{\mathbf{X}}(\mathbf{x}) = \frac{1}{(2\pi)^{p/2} |\boldsymbol{\Sigma}|^{1/2}} \exp \left(-\frac{1}{2} (\mathbf{x} - \boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1} (\mathbf{x} - \boldsymbol{\mu}) \right)$$

where $|\boldsymbol{\Sigma}|$ denotes the determinant of $\boldsymbol{\Sigma}$.

MLE: Maximum Likelihood Estimate

- Let $\mathbf{X}_1, \mathbf{X}_2, \dots, \mathbf{X}_n$ be a random sample from a multivariate normal distribution with mean vector $\boldsymbol{\mu}$ and covariance matrix $\boldsymbol{\Sigma}$.
- The log-likelihood function for the sample is given by:

$$\ell(\boldsymbol{\mu}, \boldsymbol{\Sigma}) = -\frac{n}{2} \log(2\pi) - \frac{n}{2} \log |\boldsymbol{\Sigma}| - \frac{1}{2} \sum_{i=1}^n (\mathbf{X}_i - \boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1} (\mathbf{X}_i - \boldsymbol{\mu})$$

- The MLE of $\boldsymbol{\mu}$ is the sample mean $\bar{\mathbf{X}} = \frac{1}{n} \sum_{i=1}^n \mathbf{X}_i$.

MLE: Maximum Likelihood Estimate (continued)

- To derive the MLE of Σ , we first take the derivative of the log-likelihood function with respect to Σ and set it equal to zero:

$$\frac{\partial \ell}{\partial \Sigma} = -\frac{n}{2} \Sigma^{-1} + \frac{1}{2} \sum_{i=1}^n (\mathbf{x}_i - \mu)(\mathbf{x}_i - \mu)^T \Sigma^{-2} = 0$$

- Solving for Σ , we obtain the MLE as:

$$\hat{\Sigma} = \frac{1}{n} \sum_{i=1}^n (\mathbf{x}_i - \hat{\mu})(\mathbf{x}_i - \hat{\mu})^T$$

- where $\hat{\mu}$ is the MLE of μ , as previously derived.