

Multivariate Analysis Lecture 11: Applications of PCA

Zhaoxia Yu

Professor, Department of Statistics

2025-05-05

Review of PCA

2nd PC

- **Second Principal Component.** Mathematically, we are looking for a s.t.

$$a = \arg \max_{a^T a = 1, a^T \gamma_1 = 0} a^T \Sigma a$$

- The second PC is

$$Y_2 = \gamma_2^T \mathbf{x}$$

Example: 1st PC and 2nd PC

```
gamma1=eigen(Sigma3)$vectors[,1]
gamma1
```

```
## [1] -0.8660254 -0.5000000
```

```
gamma2=eigen(Sigma3)$vectors[,2]
gamma2
```

```
## [1] 0.5000000 -0.8660254
```

- 1st PC: $-0.8660254x - 0.5y$, in the direction of γ_1 .
- 2nd PC: $0.5x - 0.8660254y$, in the direction of γ_2 .

Example: Project An Observation to 1st PC

- Consider a , the first observation in X_3 :

```
a=as.matrix(X3[1,],1,2)
a
```

```
##           x           y
## 1  1.20705  2.447848
```

- We project it to the direction of

$$\gamma_1 = (-0.8660254, -0.5)^T.$$

Because γ_1 is a unit vector

Example: Project An Observation to 1st PC

```
a
```

```
##           x           y
## 1  1.20705  2.447848
```

```
gamma1
```

```
## [1] -0.8660254 -0.5000000
```

```
a**gamma1 #inner product
```


```
##           [,1]
## 1 -2.26926
```

```
a[1]*gamma1[1]+a[2]*gamma1[2] #inner product
```

```
## [1] -2.26926
```

```
(a**gamma1)**gamma1 #projection
```

Example: Project An Observation to 1st PC



Example: Project All Observations to 1st PC

- The data matrix X_3 is a 1000×2 matrix, where each row is an observation.
- The projected values of all observations to the first PC is given by

$$Y_1 = \mathbf{X}_3 \gamma_1,$$


where

- X_3 is 1000-by-2
- γ_1 is 2-by-1, the first eigenvector of the covariance matrix Σ .

Example: Project All Observations to 1st PC

```
par(mfrow=c(1,1),pty="s")
plot(X3, xlim=c(-5,5), ylim=c(-5,5), col=8, pch=".");
abline(0,1/sqrt(3), col="red"); abline(0, -sqrt(3), col="red")
points(x=0,y=0, col="red")
proj=as.matrix(X3)%*%gamma1 %*%gamma1
text(x=proj[,1],y=proj[,2], labels="|", col="blue", srt=30, cex=0.5)
```


Example: Project All Observations to 1st PC



Example: Project All Observations to 2nd PC

```
par(mfrow=c(1,1),pty="s")
plot(X3, xlim=c(-5,5), ylim=c(-5,5), col=8, pch=".");
abline(0,1/sqrt(3), col="red"); abline(0, -sqrt(3), col="red")
points(x=0,y=0, col="red")
proj=as.matrix(X3)%*%gamma2 %*%gamma2
text(x=proj[,1],y=proj[,2], labels="|", col="blue", srt=30, cex=0.5)
```

Example: Project All Observations to 2nd PC



Example: The Two PCs

- The projected values on the first and second PCs are given by

$$PC_1 = \mathbf{X}_3 \gamma_1,$$

$$PC_2 = \mathbf{X}_3 \gamma_2$$

- Note, the above are values on the two directions, not the projections; in other words, the directional information is not included.

Example: The Two PCs



```

par(mfrow=c(1,2),pty="s")
plot(X3, xlim=c(-5,5), ylim=c(-5,5), col=8, pch=".", main="original");
abline(0,1/sqrt(3), col="red"); abline(0, -sqrt(3), col="red")
points(x=0,y=0, col="red")
projected_1=as.matrix(X3)%*%gamma1
projected_2=as.matrix(X3)%*%gamma2

plot(projected_1, projected_2, xlim=c(-5,5), ylim=c(-5,5), col=8, pch="
abline(h=0, v=0, col="red")

```

Example: The Two PCs



Example: Centered Data

- In the example, the data were generated from mean 0.
- In practice, the data may not be centered. Although the center information is not essential for PCA, it is often useful to center the data before applying PCA.
- Centering the data is equivalent to subtracting the mean from each observation.

Example: Centered Data

```



X3_centered=scale(X3, scale=FALSE)
projected_1.centered=as.matrix(X3_centered)%*%gamma1
projected_2.centered=as.matrix(X3_centered)%*%gamma2

par(mfrow=c(1,2),pty="s")
plot(X3, xlim=c(-5,5), ylim=c(-5,5), col=8, pch=".", main="original");
abline(0,1/sqrt(3), col="red"); abline(0, -sqrt(3), col="red")
points(x=0,y=0, col="red")

plot(projected_1.centered, projected_2.centered, xlim=c(-5,5), ylim=c(-
abline(h=0, v=0, col="red")

```

Example: Centered Data



Section 2

Variance Explained

Dimensionality Reduction Using PCA

- PCA aims to reduce the dimensionality of a dataset while preserving as much variance as possible.
- Dimensionality reduction using PCA can
 - ease visualization and analysis
 - help identify underlying patterns or structure in the data.
 - improve the performance of machine learning algorithms by reducing noise and collinearity.
- PCA is most effective when the majority of the variance can be captured by a small number of principal components.

Total Variance

- There are at least two justifications to use $\sum_{i=1}^p \lambda_i$ as the total variance:
 - ① One way to quantify the total variance in \mathbf{X} is the trace of the covariance matrix $\mathbf{\Sigma}$
 - Recall that $tr(\mathbf{\Sigma}) = \sum_{i=1}^n \lambda_i$
 - ② The variance of i th PC is $Var(Y_i) = Var(\gamma_i^T \mathbf{X}) = \lambda_i$. Thus, the total variance of PCs is

$$\sum_{i=1}^p Var(Y_i) = \sum_{i=1}^p \lambda_i$$

Variance Explained

- The proportion of the variance explained by i th PC explains is

$$\lambda_i / \sum_{i=1}^p \lambda_i$$

- Cumulative Explained Variance is the proportion of the total variance explained by the first k PCs is

$$\frac{\sum_{i=1}^k \lambda_i}{\sum_{i=1}^p \lambda_i}$$

Variance Explained: Example

```
lambda=eigen(Sigma3)$values
# Prop of Variance explained by 1st PC
lambda[1]/sum(lambda)
```

```
## [1] 0.8
```

```
# Prop of Cumulative Variance explained by two PCs
sum(lambda[1:2])/sum(lambda)
```

```
## [1] 1
```

- In this example, 1st PC explains 80% of variance, 2nd PC explains 20% of variance.

Section 3

Apply PCA to Data

Estimate Σ .

- We have discussed how to find PCs for a random vector $\mathbf{X}_{p \times 1} \sim (0, \Sigma)$, where Σ is a known covariance matrix.
- In practice,
 - the observed data is an $n \times p$ data matrix $\mathbf{X}_{n \times p}$
 - Σ is unknown, which can be estimated by the sample covariance matrix \mathbf{S} .

Apply PCA to A Data Set: Steps

- 1 Estimate Σ by the sample covariance matrix \mathbf{S}
- 2 Compute the eigenvectors of \mathbf{S} , and denote them by $\gamma_1, \dots, \gamma_p$
- 3 Compute the PCs
 - PC1: $Y_1 = \mathbf{X}\gamma_1$, which is a $n \times 1$ vector
 - PC2: $Y_2 = \mathbf{X}\gamma_2$, which is a $n \times 1$ vector
 -
 - PCi: $Y_i = \mathbf{X}\gamma_i$, which is a $n \times 1$ vector
 - Equivalently, compute

$$Y_{pc} = \mathbf{X}\Gamma$$

which is the $n \times p$ matrix with the i th column being the i th PC.

Section 4

Example: Iris Data

Outline

- PCA using raw data
- PCA using centered data
- PCA using standardized data
- Scree plot
- Biplot

Subsection 1

PCA Using Raw Data

Estimate Σ

```
#####
## iris data
##rearrange the data such as the response matrix is an n-by-p matrix
Y=cbind(SepalL=c(iris3[,1,1],iris3[,1,2],iris3[,1,3]),
SepalW=c(iris3[,2,1],iris3[,2,2],iris3[,2,3]),
PetalL=c(iris3[,3,1],iris3[,3,2],iris3[,3,3]),
PetalW=c(iris3[,4,1],iris3[,4,2],iris3[,4,3]))
#for unknown reasons, data.frame won't work but cbind works
#alternatively, we can use the following way to define y
#y=aperm(iris3,c(1,3,2));dim(y)=c(150,4)
S=cov(Y)
```

Compute PCs

```
eigen.vec=eigen(S)$vectors
eigen.val=eigen(S)$values
gamma1=eigen.vec[,1]
gamma2=eigen.vec[,2]
gamma3=eigen.vec[,3]
gamma4=eigen.vec[,4]
# The four pcs:
pc1=Y%*%gamma1
pc2=Y%*%gamma2
pc3=Y%*%gamma3
pc4=Y%*%gamma4
# Equivalently, we can obtain the nx4 matrix of PCs
PCs=Y%*%eigen.vec
```

The Loadings (Weights)

```
gamma1
```

```
## [1] 0.36138659 -0.08452251 0.85667061 0.35828920
```

```
gamma2
```

```
## [1] -0.65658877 -0.73016143 0.17337266 0.07548102
```

```
gamma3
```

```
## [1] -0.58202985 0.59791083 0.07623608 0.54583143
```

```
gamma4
```


```
## [1] 0.3154872 -0.3197231 -0.4798390 0.7536574
```

The Four PCs

- PC 1: $Y_1 = 0.36SL - 0.08SW + 0.86PL + 0.36PL$
- PC 2: $Y_2 = -0.66SL - 0.73SW + 0.17PL + 0.08PL$
- PC 3: $Y_3 = -0.60SL - 0.60SW + 0.08PL + 0.55PL$
- PC 4: $Y_4 = 0.32SL - 0.32SW - 0.48PL + 0.75PL$
- Note that PC1 loads the most on PL. This is not surprising because PL has the largest variance among all the four features.

Visualize PC1 and PC2

```
plot(pc1,pc2,xlab="PC1", ylab="PC2", type="n")
points(pc1[1:50], pc2[1:50], col=1, pch="s")
points(pc1[51:100], pc2[51:100], col=2, pch="e")
points(pc1[101:150], pc2[101:150], col=3, pch="i")
```




Proportions of Variation Explained (Cumulative)

```
cumsum(eigen.val/sum(eigen.val))
```

```
## [1] 0.9246187 0.9776852 0.9947878 1.0000000
```

```
plot(cumsum(eigen.val)/sum(eigen.val), type="b", main="Iris Data: Variation Explained",  
ylab="eigenvalues")
```



Subsection 2

PCA Using Centered Data

PCA Using Centered Data

The four pcs:

```
pc1_c=scale(Y, scale=FALSE)%*%gamma1
```

```
pc2_c=scale(Y, scale=FALSE)%*%gamma2
```

```
pc3_c=scale(Y, scale=FALSE)%*%gamma3
```

```
pc4_c=scale(Y, scale=FALSE)%*%gamma4
```

or

```
PCs_c=scale(Y, scale=FALSE)%*%eigen(S)$vectors
```

The PCs from R


```
#help(princomp), pay attention to fix_sign
obj=princomp(Y)
names(obj)
```

```
## [1] "sdev"      "loadings" "center"    "scale"     "n.obs"     "scores"
```

```
#help(loadings), pay attention to cutoff
#check loadings and scores to verify that
#they are same as what we calculated
```

The PCs from R vs our PCs


```
plot(obj$scores[,1], pc1)
```



- The difference is due to not centering vs centering

R Uses Centered Data to Compute PCA

```
plot(obj$scores[,1], pc1_c)
```



- They are identical

Subsection 3

PCA Using Standardized Data

Covariance of Standardized Data

- The covariance matrix using standardized data

```
cov(scale(Y, scale=TRUE))
```

```
##           SepalL      SepalW      PetalL      PetalW
## SepalL  1.0000000 -0.1175698  0.8717538  0.8179411
## SepalW -0.1175698  1.0000000 -0.4284401 -0.3661259
## PetalL  0.8717538 -0.4284401  1.0000000  0.9628654
## PetalW  0.8179411 -0.3661259  0.9628654  1.0000000
```


Correlation Matrix

- The correlation matrix is the covariance matrix of standardized data

```
cor(Y)
```


```
##          SepalL      SepalW      PetalL      PetalW
## SepalL  1.0000000 -0.1175698  0.8717538  0.8179411
## SepalW -0.1175698  1.0000000 -0.4284401 -0.3661259
## PetalL  0.8717538 -0.4284401  1.0000000  0.9628654
## PetalW  0.8179411 -0.3661259  0.9628654  1.0000000
```

PCA Based on Correlation Matrix

```
gamma1_s=eigen(cor(Y))$vectors[,1]  
gamma2_s=eigen(cor(Y))$vectors[,2]  
gamma3_s=eigen(cor(Y))$vectors[,3]  
gamma4_s=eigen(cor(Y))$vectors[,4]
```

The PCs based on Correlation Matrix

```
PCs_s=scale(Y)%*% eigen(cor(Y))$vectors
plot(PCs_s[,1:2], xlab="PC1", ylab="PC2",
     main="Iris Data: 1st and 2nd PC (Corr)", type="n")
points(PCs_s[1:50, 1:2], col=1, pch="s")
points(PCs_s[51:100, 1:2], col=2, pch="e")
points(PCs_s[101:150, 1:2], col=3, pch="i")
```



PCA with Standardized Data in R

- The princomp function in R has an option to use correlation matrix rather than covariance matrix

```
obj.cor=princomp(Y, cor=TRUE)
```

Standardized vs Not-Standardized

• Loadings

```
obj$loadings[,1:4]
```

```
##              Comp.1      Comp.2      Comp.3      Comp.4
## SepalL  0.36138659  0.65658877  0.58202985  0.3154872
## SepalW -0.08452251  0.73016143 -0.59791083 -0.3197231
## PetalL  0.85667061 -0.17337266 -0.07623608 -0.4798390
## PetalW  0.35828920 -0.07548102 -0.54583143  0.7536574
```

```
obj.cor$loading[,1:4]
```

```
##              Comp.1      Comp.2      Comp.3      Comp.4
## SepalL  0.5210659  0.37741762  0.7195664  0.2612863
## SepalW -0.2693474  0.92329566 -0.2443818 -0.1235096
## PetalL  0.5804131  0.02449161 -0.1421264 -0.8014492
## PetalW  0.5648565  0.06694199 -0.6342727  0.5235971
```

PCA based on Cov vs Cor: 1st and 2nd PC


```

par(mfrow=c(1,2))
plot(PCs_c[,1:2], main="PCA using Cov", type="n")
points(PCs_c[1:50, 1:2], col=1, pch="s")
points(PCs_c[51:100, 1:2], col=2, pch="e")
points(PCs_c[101:150, 1:2], col=3, pch="i")
plot(PCs_s[,1:2], main="PCA using Cor", type="n")
points(PCs_s[1:50, 1:2], col=1, pch="s")
points(PCs_s[51:100, 1:2], col=2, pch="e")
points(PCs_s[101:150, 1:2], col=3, pch="i")


```

PCA Using Standardized Data

PCA using Cov



PCA using Cor



PCA based on Cov vs Cor: Variance Explained (Cumulatively)

```



par(mfrow=c(1,2))
cumsum(eigen(cov(Y))$values)/sum(eigen(cov(Y))$values)
plot(cumsum(eigen(cov(Y))$values)/sum(eigen(cov(Y))$values),
     type="b", main="Iris Data: Cov",ylab="Cum Variance Explained")
cumsum(eigen(cor(Y))$values)/sum(eigen(cor(Y))$values)
plot(cumsum(eigen(cor(Y))$values)/sum(eigen(cor(Y))$values),
     type="b", main="Iris Data: Cor",ylab="Cum Variance Explained")

```

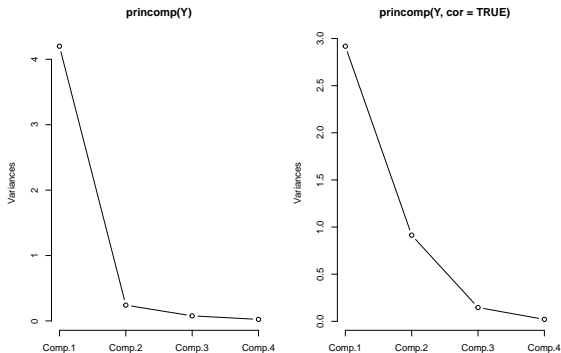

PCA Using Standardized Data

```
## [1] 0.9246187 0.9776852 0.9947878 1.0000000
```

```
## [1] 0.7296245 0.9581321 0.9948213 1.0000000
```




```
par(mfrow=c(1,2))
plot(princomp(Y), type="l")
plot(princomp(Y, cor=TRUE), type="l")
```



PCA based on Cov vs Cor: Variances (Scree Plot)

```
par(mfrow=c(1,2))
plot(princomp(Y))
plot(princomp(Y, cor=TRUE))
```



Visualizing PCA

- We have seen scree plots for visualizing the variance explained by each PC
- Biplots are another way to visualize PCA results
- A biplot is a visualization tool that combines:
 - Scores (projected data points)
 - Loadings (variable vectors) from Principal Component Analysis (PCA).
- A biplot shows relationships between observations and between variables in reduced dimensions (e.g., PC1 vs. PC2).

Biplot: loadings

```
obj$loadings
```

```
##
```

```
## Loadings:
```

```
##          Comp.1 Comp.2 Comp.3 Comp.4
```

```
## SepalL  0.361  0.657  0.582  0.315
```

```
## SepalW           0.730 -0.598 -0.320
```

```
## PetalL  0.857 -0.173           -0.480
```

```
## PetalW  0.358           -0.546  0.754
```

```
##
```

```
##          Comp.1 Comp.2 Comp.3 Comp.4
```

```
## SS loadings      1.00   1.00   1.00   1.00
```


```
## Proportion Var   0.25   0.25   0.25   0.25
```

```
## Cumulative Var   0.25   0.50   0.75   1.00
```

PCA Using Standardized Data

Biplot

```
biplot(obj, main="Iris Data: Biplot")
```



Biplot: the scale parameter

- The loadings and scores often have very different scales.
- The scale parameter harmonizes them for visualization.
- The default value is 1, which means the loadings are scaled by 1 and the scores are scaled by 1, i.e., cores and loadings are scaled equally.
- Make scale < 1 if variable arrows are too small to interpret.
- Use scale > 1 if data points are too crowded.

Section 5

Example: Wine Data

The Wine Data

```
wine = read.csv("http://archive.ics.uci.edu/ml/machine-learning-databases/wine/wine.data")
dim(wine)
```

```
## [1] 4898 12
```

```
head(wine)
```



```
##   fixed.acidity volatile.acidity citric.acid residual.sugar chloride
## 1           7.0           0.27           0.36           20.7       0.04
## 2           6.3           0.30           0.34           1.6       0.04
## 3           8.1           0.28           0.40           6.9       0.05
## 4           7.2           0.23           0.32           8.5       0.05
## 5           7.2           0.23           0.32           8.5       0.05
## 6           8.1           0.28           0.40           6.9       0.05
##   free.sulfur.dioxide total.sulfur.dioxide density    pH sulphates al
## 1                   45                   170 1.0010 3.00       0.45
## 2                   14                   132 0.9940 3.30       0.49
## 3                   30                   97  0.9951 3.26       0.44
```

PCA based on Cov vs Cor: Variance Explained

```
par(mfrow=c(1,2))
cumsum(eigen(cov(wine))$values)/sum(eigen(cov(wine))$values)
plot(cumsum(eigen(cov(wine))$values)/sum(eigen(cov(wine))$values),
     type="b", main="Wine Data: Cov",ylab="Cum Variance Explained")
cumsum(eigen(cor(wine))$values)/sum(eigen(cor(wine))$values)
plot(cumsum(eigen(cor(wine))$values)/sum(eigen(cor(wine))$values),
     type="b", main="Wine Data: Cor",ylab="Cum Variance Explained")
```



```
## [1] 0.9093312 0.9886453 0.9987968 0.9994164 0.9997397 0.9999753 0.9
## [8] 0.9999906 0.9999960 0.9999998 1.0000000 1.0000000

## [1] 0.2788891 0.4110633 0.5253276 0.6157327 0.6970063 0.7752366 0.8
## [8] 0.8996651 0.9460402 0.9746180 0.9982942 1.0000000
```





PCA based on Cov vs Cor: Variances (Scree Plot)

```
par(mfrow=c(1,2))
plot(princomp(wine), type="l")
plot(princomp(wine, cor=TRUE), type="l")
```





PCA based on Cov vs Cor: Variances (Scree Plot)

```
par(mfrow=c(1,2))
plot(princomp(wine))
plot(princomp(wine, cor=TRUE))
```



PCA based on Cov vs Cor: Loadings of 1st PC

```
par(mfrow=c(1,2))  
plot(princomp(wine)$loadings[,1], ylab="loadings of 1st PC")  
plot(princomp(wine, cor=TRUE)$loadings[,1], ylab="loadings of 1st PC")
```



Wine Data: PCA_COV and PCA_COR Are Different

- In the wine data, raw data and standardized data give very different results
- Raw Data: the first PC dominates the rest
- Standardized data: the variances of the PCs are less different
- Why?

Wine Data: The measurements

- There are 12 measurements / features in the wine data
- The variances are very different

```
apply(wine, 2, var)
```

##	fixed.acidity	volatile.acidity	citric.acid
##	7.121136e-01	1.015954e-02	1.464579e-02
##	residual.sugar	chlorides	free.sulfur.dioxide
##	2.572577e+01	4.773337e-04	2.892427e+02
##	total.sulfur.dioxide	density	pH
##	1.806085e+03	8.945524e-06	2.280118e-02
##	sulphates	alcohol	quality
##	1.302471e-02	1.514427e+00	7.843557e-01

Standardized vs Not-Standardized

- PCA depends on the measurement scale
- When features are not standardized, the leading PCs tend to give larger loadings for features with large variances
- When similar variables are measured using different units, standardization is typically recommended before PCA

Section 6

Choose k

Choose the Number of PCs

- A rule of thumb? Choose the PC's that contain more information than the average amount of information per PC.
- People often choose k such that at least a certain percentage of the total variance is explained. E.g., 80%, 90%, 95%, 99%
- Scree Plots:
 - A scree plot is a plot of the eigenvalues against the number of PCs
 - we look for an “elbow point” where the decrease in eigenvalues becomes less steep. The number of PCs corresponding to the elbow point can be chosen.
- No unified solution. The choice of the number of PCs may depend on the specific problem

Choose the Number of PCs

- Balancing dimensionality reduction and information retention.
- Too few PCs may result in significant loss of information
- Too many PCs is not efficient in dimension reduction application or research question

The Remaining Lectures Will Cover

- Lec 12: Linear Discriminant Analysis
- Lec 13: Linear Discriminant Analysis
- Lec 14: Factor Analysis
- Lec 15: Cluster Analysis
- Lec 16: Canonical Analysis
- Lec 17: Structural Equation Modeling
- Lec 18: Conclusion